

NEEL BHANDARI

neelbhandari6.github.io | neelbhandari64@gmail.com | linkedin.com/in/neelbhandari | +1-412-284-8313

EDUCATION

Carnegie Mellon University

M.S. in Machine Learning and NLP (MIIS), Language Technologies Institute

Pittsburgh, PA

Aug. 2024 - Dec. 2025

Coursework (Ongoing): Introduction to Machine Learning, Advanced NLP, On-Device Machine Learning.

R.V. College of Engineering

B.Tech in Computer Science and Engineering, GPA: 8.63/10

Bangalore, India

Aug. 2018-Sept. 2022

WORK EXPERIENCE

Machine Learning Scientist

Aug 2022- July 2024

PayPal Inc.

Bangalore, India

- Spearheaded research and development behind the first automated Root Cause Identification tool at PayPal, which promised to reduce root cause analysis time by over **75%**.
- Developed, trained and tested language models on proprietary data along with use-case specific metrics for the first LLM-based PayPal Assistant chatbot.
- Led experimentation and testing through fine-tuning and prompt engineering to develop a customer-centric personalisation feature, aimed at solving account specific questions for consumers.

Machine Learning Intern

April 2022 - July 2022

PayPal Inc.

Bangalore, India

- Developed a proof-of-concept Root Cause Identification tool focused on automated extraction of problem statements and root causes from unstructured chat transcripts.
- Experimented and integrated several techniques such as question-answering, summarisation and cluster-based modelling to create a robust pipeline.

RESEARCH PROJECTS

Investigating and Mitigating Hallucinations in RAG systems

Aug. 2020 – June 2021

Graduate Research Student, Advised by Prof. Maarten Sap

Pittsburgh, PA

- Developing a systematic approach to analyze origin points and effects of hallucinations across Retrieval-Augmented Generation (RAG) system components, with a specific focus on various types of knowledge conflicts.
- Formulating synthetic data generation methods to enhance robustness in RAG pipelines.

Aya Open Science Project [1]

Jan. 2023 - April 2024

Member of Core Research Team, Cohere For AI Community

Remote

- Led experimentation and design on quality estimation and human-led pruning of multilingual datasets to understand the impact of data quality on model performance.
- Engineered the development of the data processing pipeline for diverse multilingual datasets, reducing the caching time from several days down to hours.

Robustness of Text-Based Adversarial Attacks [2] [3]

June 2021 - Mar. 2022

External Research Student, MIT-IBM Watson AI Lab

Remote

- Collaborated with Dr. Pin-Yu Chen on studying robustness of text-based adversarial attacks from a multilingual lens
- Designed and developed a novel attack-agnostic algorithm that showcased significant improvement in robustness over existing attacks in a round-trip translation setting.

SELECTED PUBLICATIONS

1. Ahmet Üstün, Viraat Aryabumi, ... , Neel Bhandari et. al. **Aya Model: An Instruction Finetuned Open-Access Multilingual Language Model**, *ACL 2024 - Best Paper*
2. Neel Bhandari, Pin-Yu Chen. **Lost In Translation: Generating Adversarial Examples Robust to Round-Trip Translation (Full Paper)**, *ICASSP 2023*.
3. Neel Bhandari, Pin-Yu Chen. **Lost In Translation: Generating Adversarial Examples Robust to Round-Trip Translation**, *Workshop on Socially Responsible Machine Learning, ICLR, 2022*.

SKILLS

- **Programming Languages & Libraries:** Python; C++; SQL; NumPy; Pandas; PyTorch; Tensorflow; Scikit-learn